

# model\_03a\_bert\_base\_cos\_sim\_LR\_insight

May 3, 2023

## 1 Model 03a

Evidence retrieval using a Siamese BERT classification model. This is similar to Model 01, however, it only uses official pre-trained models from hugging face.

Ref: - [Hugging face pre-trained models](#) - [Hugging face guide to fine-tuning](#) - [Hugging face guide to fine-tuning easy](#) - [SO Guide](#)

### 1.1 Setup

#### 1.1.1 Working Directory

```
[ ]: # Change the working directory to project root
from pathlib import Path
import os
ROOT_DIR = Path.cwd()
while not ROOT_DIR.joinpath("src").exists():
    ROOT_DIR = ROOT_DIR.parent
os.chdir(ROOT_DIR)
```

#### 1.1.2 File paths

```
[ ]: MODEL_PATH = ROOT_DIR.joinpath("./result/models/*")
DATA_PATH = ROOT_DIR.joinpath("./data/*")
NER_PATH = ROOT_DIR.joinpath("./result/ner/*")
```

#### 1.1.3 Dependencies

```
[ ]: # Imports and dependencies
import torch
from torch.utils.data import Dataset, DataLoader
from torch.nn import Module, CosineEmbeddingLoss
from transformers import BertModel, BertTokenizer
from torch.optim import Adam
from torch.optim.lr_scheduler import LinearLR
from torcheval.metrics import BinaryAccuracy, BinaryF1Score

from src.torch_utils import get_torch_device
```

```

import json
from dataclasses import dataclass
from typing import List, Union, Tuple
from tqdm import tqdm
import random
import numpy as np
from datetime import datetime
from math import exp

TORCH_DEVICE = get_torch_device()

```

/opt/homebrew/Caskroom/miniconda/base/envs/comp90042\_project/lib/python3.8/site-packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See [https://ipywidgets.readthedocs.io/en/stable/user\\_install.html](https://ipywidgets.readthedocs.io/en/stable/user_install.html)

```

from .autonotebook import tqdm as notebook_tqdm

```

Torch device is 'mps'

## 1.2 Dataset

```

[ ]: @dataclass
class ClaimEvidencePair:
    claim_id:str
    evidence_id:str
    label:int = 0

```

```

[ ]: class SiameseEvalDataset(Dataset):
    def __init__(
        self,
        dev_claims_path:Path,
        evidence_path:Path,
        device = None,
        verbose:bool=True
    ) -> None:
        super(SiameseEvalDataset, self).__init__()
        self.verbose = verbose
        self.device = device

        # Load claims data from json
        with open(dev_claims_path, mode="r") as f:
            self.claims = (json.load(fp=f))

        # Load evidence library
        self.evidence = dict()
        with open(evidence_path, mode="r") as f:
            self.evidence.update(json.load(fp=f))

```

```

        # Get a list of all evidences within the dev set
        self.related_evidences = sorted({
            evidence_id
            for claim in self.claims.values()
            for evidence_id in claim["evidences"]
        })

        # Generate the data
        self.data = self.__generate_data()
        return

    def __generate_data(self):
        data = []
        for claim_id, claim in tqdm(
            iterable=self.claims.items(),
            desc="claims",
            disable=not self.verbose
        ):
            evidence_ids = claim["evidences"]

            # Get the positives
            for evidence_id in evidence_ids:
                data.append(ClaimEvidencePair(
                    claim_id=claim_id,
                    evidence_id=evidence_id,
                    label=1
                ))

            # Get some negatives
            n_neg = 0
            for rel_evidence_id in self.related_evidences:
                if n_neg >= 10:
                    break
                if rel_evidence_id in evidence_ids:
                    continue
                data.append(ClaimEvidencePair(
                    claim_id=claim_id,
                    evidence_id=rel_evidence_id,
                    label=-1
                ))
                n_neg += 1
            return data

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx) -> Tuple[Union[str, torch.Tensor]]:

```

```

# Fetch the required data rows
data = self.data[idx]

# Get the label
label = torch.tensor(data.label, device=self.device)

# Get text ids
claim_id = data.claim_id
evidence_id = data.evidence_id

# Get text
claim_text = self.claims[claim_id]["claim_text"]
evidence_text = self.evidence[evidence_id]

return (claim_text, evidence_text, label)

```

```

[ ]: class SiameseDataset(Dataset):

    def __init__(
        self,
        claims_paths:List[Path],
        claims_shortlist_paths:List[Path],
        evidence_path:Path,
        evidence_shortlists:List[Path] = None,
        device = None,
        n_neg_shortlist:int = 10,
        n_neg_general:int = 10,
        verbose:bool=True
    ) -> None:
        super(SiameseDataset, self).__init__()
        self.verbose = verbose
        self.device = device
        self.n_neg_shortlist = n_neg_shortlist
        self.n_neg_general = n_neg_general

        # Load claims data from json, this is a list as we could use
        # multiple json files in the same dataset
        self.claims = dict()
        for json_file in claims_paths:
            with open(json_file, mode="r") as f:
                self.claims.update(json.load(fp=f))
                # print(f"loaded claims: {json_file}")

        # Load the pre-retrieved shortlist of evidences by claim
        self.claims_shortlist = dict()
        for json_file in claims_shortlist_paths:
            with open(json_file, mode="r") as f:

```

```

        self.claims_shortlist.update(json.load(fp=f))
        # print(f"loaded claims_shortlist: {json_file}")

    # Load evidence library
    self.evidence = dict()
    with open(evidence_path, mode="r") as f:
        self.evidence.update(json.load(fp=f))
        # print(f"loaded evidences: {json_file}")

    # Load the evidence shortlists if available
    # Reduce the overall evidence list to the shortlist
    if evidence_shortlists is not None:
        self.evidence_shortlist = set()
        for json_file in evidence_shortlists:
            with open(json_file, mode="r") as f:
                self.evidence_shortlist.update(json.load(fp=f))
                # print(f"loaded evidence shortlist: {json_file}")

    # print(f"n_evidences: {len(self.evidence)}")

    # Generate the data
    self.data = self.__generate_data()
    return

def __generate_data(self):
    print("Generate siamese dataset")

    data = []
    for claim_id, claim in tqdm(
        iterable=self.claims.items(),
        desc="claims",
        disable=not self.verbose
    ):
        # Check if we have evidences supplied, this will inform
        # whether this is for training
        is_training = "evidences" in claim.keys()
        pos_evidence_ids = set()

        # Get positive samples from evidences with label=1
        if is_training:
            pos_evidence_ids.update(claim["evidences"])

        for evidence_id in pos_evidence_ids:
            data.append(ClaimEvidencePair(
                claim_id=claim_id,
                evidence_id=evidence_id,
                label=1
            ))

```

```

    ))

    # Get negative samples from pre-retrieved evidences
    # for each claim with label=-1
    retrieved_evidence_ids = self.claims_shortlist.get(claim_id, [])
    if len(retrieved_evidence_ids) > 0:
        retrieved_neg_evidence_ids = random.sample(
            population=retrieved_evidence_ids,
            k=min(self.n_neg_shortlist, len(retrieved_evidence_ids))
        )

        # Generate claim and shortlisted negative evidence pairs
        for evidence_id in retrieved_neg_evidence_ids:
            data.append(ClaimEvidencePair(
                claim_id=claim_id,
                evidence_id=evidence_id,
                label=-1
            ))

    # Get negative samples from shortlisted evidences list with label=0
    if len(self.evidence_shortlist) > 0:
        shortlist_neg_evidence_ids = random.sample(
            population=self.evidence_shortlist,
            k=min(self.n_neg_general, len(self.evidence_shortlist))
        )

        # Generate claim and shortlisted negative evidence pairs
        for evidence_id in shortlist_neg_evidence_ids:
            data.append(ClaimEvidencePair(
                claim_id=claim_id,
                evidence_id=evidence_id,
                label=-1
            ))

    continue

print(f"Generated data n={len(data)}")

return data

def __len__(self):
    return len(self.data)

def __getitem__(self, idx) -> Tuple[Union[str, torch.Tensor]]:
    # Fetch the required data rows
    data = self.data[idx]

```

```

    # Get the label
    label = torch.tensor(data.label, device=self.device)

    # Get text ids
    claim_id = data.claim_id
    evidence_id = data.evidence_id

    # Get text
    claim_text = self.claims[claim_id]["claim_text"]
    evidence_text = self.evidence[evidence_id]

    return (claim_text, evidence_text, label)

```

```
[ ]: # WE WILL GENERATE THE DATASET PER EPOCH SO TO RANDOMISE THE NEGATIVE SAMPLES
```

```

# train_data = SiameseDataset(
#     claims_paths=[DATA_PATH.with_name("train-claims.json")],
#     claims_shortlist_paths=[NER_PATH.
# ↪with_name("train_claim_evidence_retrieved.json")],
#     evidence_shortlists=[NER_PATH.
# ↪with_name("shortlist_train_claim_evidence_retrieved.json")],
#     evidence_path=DATA_PATH.with_name("evidence.json"),
#     device=TORCH_DEVICE,
#     n_neg_shortlist=100,
#     n_neg_general=100
# )

```

### 1.3 Build model

```
[ ]: class SiameseEmbedderBert(Module):

    def __init__(
        self,
        pretrained_name:str,
        device,
        **kwargs
    ) -> None:
        super(SiameseEmbedderBert, self).__init__(**kwargs)
        self.device = device

        # Use a pretrained tokenizer
        self.tokenizer = BertTokenizer.from_pretrained(pretrained_name)

        # Use a pretrained model
        self.bert = BertModel.from_pretrained(pretrained_name)
        self.bert.to(device=device)
        return

```

```

def forward(self, claim_texts, evidence_texts, eval_mode:bool=False) ->
↳ Tuple[torch.Tensor]:

    # Run the tokenizer
    t_kwargs = {
        "return_tensors": "pt",
        "padding": True,
        "truncation": True,
        "max_length": 100,
        "add_special_tokens": True
    }
    claim_x = self.tokenizer(claim_texts, **t_kwargs)
    evidence_x = self.tokenizer(evidence_texts, **t_kwargs)

    claim_x = claim_x["input_ids"].to(device=self.device)
    evidence_x = evidence_x["input_ids"].to(device=self.device)

    # Run Bert
    claim_x = self.bert(claim_x, return_dict=True).pooler_output
    evidence_x = self.bert(evidence_x, return_dict=True).pooler_output
    # dim=768

    # Cosine similarity
    if eval_mode:
        cos_sim = torch.cosine_similarity(x1=claim_x, x2=evidence_x)
        return claim_x, evidence_x, cos_sim

    return claim_x, evidence_x

```

## 1.4 Training and evaluation loop

```

[ ]: model = SiameseEmbedderBert(
    pretrained_name="bert-base-cased",
    device=TORCH_DEVICE
)

```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertModel: ['cls.predictions.decoder.weight',

'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.bias', 'cls.seq\_relationship.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.dense.bias', 'cls.seq\_relationship.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of



a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```
[ ]: loss_fn = CosineEmbeddingLoss()
optimizer = Adam(
    params=model.parameters(),
    lr=0.000002
) #! Hyperparams
```

```
[ ]: run_time = datetime.now().strftime('%Y_%m_%d_%H_%M')
MODEL_NAME = f"model_03a_bert_base_cos_sim_{run_time}.pth"
N_EPOCHS = 100
BATCH_SIZE = 64
```

```
[ ]: dev_data = SiameseEvalDataset(
    dev_claims_path=DATA_PATH.with_name("dev-claims.json"),
    evidence_path=DATA_PATH.with_name("evidence.json"),
    device=TORCH_DEVICE
)

dev_dataloader = DataLoader(
    dataset=dev_data,
    shuffle=False,
    batch_size=BATCH_SIZE
)
```

```
claims: 100%|          | 154/154 [00:00<00:00, 179472.86it/s]
claims: 100%|          | 154/154 [00:00<00:00, 179472.86it/s]
```

```
[ ]: import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: # Run evaluation before training to establish baseline
model.eval()

dev_batches = tqdm(dev_dataloader, desc="dev batches")
epoch_pos_cos_sim = []
epoch_neg_cos_sim = []
for batch in dev_batches:
    claim_texts, evidence_texts, labels = batch

    # Forward
    claim_emb, evidence_emb, cos_sim = model(claim_texts, evidence_texts,
    ↪eval_mode=True)

    # Cosine similarity
    labelled_cos_sim = cos_sim * labels
```

```

pos_cos_sim = labelled_cos_sim[torch.where(labelled_cos_sim > 0)]
neg_cos_sim = labelled_cos_sim[torch.where(labelled_cos_sim < 0)]

batch_pos_cos_sim = torch.mean(pos_cos_sim).cpu().item()
batch_neg_cos_sim = torch.mean(neg_cos_sim).cpu().item() * -1

epoch_pos_cos_sim.append(batch_pos_cos_sim)
epoch_neg_cos_sim.append(batch_neg_cos_sim)

dev_batches.postfix = f"pos cos_sim: {batch_pos_cos_sim:.3f}" + \
    f" neg cos_sim: {batch_neg_cos_sim:.3f}"

continue

```

```

dev batches: 100%|          | 32/32 [00:09<00:00, 3.51it/s, pos cos_sim: 0.972
neg cos_sim: 0.868]

```

```

[ ]: print(f"Average cos sim (pos, neg): {np.mean(epoch_pos_cos_sim):3f}, {np.
      ↪mean(epoch_neg_cos_sim):3f}")

```

```

Average cos sim (pos, neg): 0.863419, 0.838396

```

```

[ ]: metric_accuracy = BinaryAccuracy()
metric_f1 = BinaryF1Score()
metric_recall = BinaryF1Score()

scheduler = LinearLR(
    optimizer=optimizer,
    start_factor=0.1,
    end_factor=1,
    total_iters=int(N_EPOCHS/10),
    verbose=True
)
best_epoch_loss = 999
for epoch in range(N_EPOCHS):

    print(f"Epoch: {epoch} of {N_EPOCHS}\n")

    # Run training
    model.train()

    train_data = SiameseDataset(
        claims_paths=[DATA_PATH.with_name("train-claims.json")],
        claims_shortlist_paths=[NER_PATH.
    ↪with_name("train_claim_evidence_retrieved.json")],
        evidence_shortlists=[NER_PATH.
    ↪with_name("shortlist_train_claim_evidence_retrieved.json")],

```

```

        evidence_path=DATA_PATH.with_name("evidence.json"),
        device=TORCH_DEVICE,
        n_neg_shortlist=2,
        n_neg_general=1
    )

    train_dataloader = DataLoader(
        dataset=train_data,
        shuffle=True,
        batch_size=BATCH_SIZE
    )

    train_batches = tqdm(train_dataloader, desc="train batches")
    running_losses = []
    for batch in train_batches:
        claim_texts, evidence_texts, labels = batch

        # Reset optimizer
        optimizer.zero_grad()

        # Forward + loss
        claim_emb, evidence_emb = model(claim_texts, evidence_texts)
        loss = loss_fn(input1=claim_emb, input2=evidence_emb, target=labels)

        # Backward + optimiser
        loss.backward()
        optimizer.step()

        # Update running loss
        batch_loss = loss.item() * len(batch)
        running_losses.append(batch_loss)

        train_batches.postfix = f"loss: {batch_loss:.3f}"

        continue

    scheduler.step()

    epoch_loss = np.average(running_losses)
    print(f"Average epoch loss: {epoch_loss}")

    # Save model
    if epoch_loss <= best_epoch_loss:
        best_epoch_loss = epoch_loss
        torch.save(model, MODEL_PATH.with_name(MODEL_NAME))
        print(f"Saved model to: {MODEL_PATH.with_name(MODEL_NAME)}")

```

```

# Evaluate every 5 epochs
# if epoch % 5 != 0:
#     continue

# Run evaluation before training to establish baseline
model.eval()

dev_batches = tqdm(dev_dataloader, desc="dev batches")
epoch_pos_cos_sim = []
epoch_neg_cos_sim = []
for batch in dev_batches:
    claim_texts, evidence_texts, labels = batch

    # Forward
    claim_emb, evidence_emb, cos_sim = model(claim_texts, evidence_texts,
    ↪eval_mode=True)

    # Cosine similarity
    labelled_cos_sim = cos_sim * labels
    pos_cos_sim = labelled_cos_sim[torch.where(labelled_cos_sim > 0)]
    neg_cos_sim = labelled_cos_sim[torch.where(labelled_cos_sim < 0)]

    batch_pos_cos_sim = torch.mean(pos_cos_sim).cpu().item()
    batch_neg_cos_sim = torch.mean(neg_cos_sim).cpu().item() * -1

    epoch_pos_cos_sim.append(batch_pos_cos_sim)
    epoch_neg_cos_sim.append(batch_neg_cos_sim)

    dev_batches.postfix = f"pos cos_sim: {batch_pos_cos_sim:.3f}" + \
        f" neg cos_sim: {batch_neg_cos_sim:.3f}"

    continue

print(f"Average cos sim (pos, neg): {np.mean(epoch_pos_cos_sim):3f}, {np.
    ↪mean(epoch_neg_cos_sim):3f}")

print("Done!")

```

Adjusting learning rate of group 0 to 2.0000e-07.

Epoch: 0 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 157.84it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:48<00:00, 1.11it/s, loss: 1.960]

Adjusting learning rate of group 0 to 3.8000e-07.  
Average epoch loss: 1.41377578221865  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.25it/s, pos cos\_sim: 0.970  
neg cos\_sim: 0.888]

Average cos sim (pos, neg): 0.879555, 0.861016  
Epoch: 1 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 155.89it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:46<00:00, 1.14it/s, loss: 1.451]

Adjusting learning rate of group 0 to 5.6000e-07.  
Average epoch loss: 1.3904771886088632  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.29it/s, pos cos\_sim: 0.947  
neg cos\_sim: 0.851]

Average cos sim (pos, neg): 0.837746, 0.812627  
Epoch: 2 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 162.05it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:46<00:00, 1.14it/s, loss: 1.118]

Adjusting learning rate of group 0 to 7.4000e-07.  
Average epoch loss: 1.3746617729506216  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.29it/s, pos cos\_sim: 0.910  
neg cos\_sim: 0.819]

Average cos sim (pos, neg): 0.825345, 0.786491  
Epoch: 3 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 168.03it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:46<00:00, 1.14it/s, loss: 0.821]

Adjusting learning rate of group 0 to 9.2000e-07.  
Average epoch loss: 1.3272475949988878  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.29it/s, pos cos\_sim: 0.896  
neg cos\_sim: 0.878]

Average cos sim (pos, neg): 0.865243, 0.848583  
Epoch: 4 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 172.60it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:45<00:00, 1.14it/s, loss: 0.392]

Adjusting learning rate of group 0 to 1.1000e-06.

Average epoch loss: 0.9673865113622886

Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.28it/s, pos cos\_sim: 0.926  
neg cos\_sim: 0.842]

Average cos sim (pos, neg): 0.835551, 0.822032  
Epoch: 5 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 159.28it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:46<00:00, 1.14it/s, loss: 1.265]

Adjusting learning rate of group 0 to 1.2800e-06.

Average epoch loss: 0.7218735043174964

Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.30it/s, pos cos\_sim: 0.942  
neg cos\_sim: 0.882]

Average cos sim (pos, neg): 0.901280, 0.867277  
Epoch: 6 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 173.51it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:46<00:00, 1.14it/s, loss: 0.648]

Adjusting learning rate of group 0 to 1.4600e-06.  
Average epoch loss: 0.6569000715928629  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.29it/s, pos cos\_sim: 0.924  
neg cos\_sim: 0.867]

Average cos sim (pos, neg): 0.888637, 0.858377  
Epoch: 7 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 164.89it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:45<00:00, 1.15it/s, loss: 1.187]

Adjusting learning rate of group 0 to 1.6400e-06.  
Average epoch loss: 0.6235487781280329  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.25it/s, pos cos\_sim: 0.832  
neg cos\_sim: 0.840]

Average cos sim (pos, neg): 0.823294, 0.844328  
Epoch: 8 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 165.54it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:46<00:00, 1.14it/s, loss: 0.087]

Adjusting learning rate of group 0 to 1.8200e-06.  
Average epoch loss: 0.596648235074129  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.27it/s, pos cos\_sim: 0.939  
neg cos\_sim: 0.837]

Average cos sim (pos, neg): 0.903299, 0.826131  
Epoch: 9 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 166.75it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:45<00:00, 1.14it/s, loss: 0.613]

Adjusting learning rate of group 0 to 2.0000e-06.  
Average epoch loss: 0.587202386730466  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.24it/s, pos cos\_sim: 0.622  
neg cos\_sim: 0.891]

Average cos sim (pos, neg): 0.764762, 0.838932  
Epoch: 10 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 165.99it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:45<00:00, 1.15it/s, loss: 0.801]

Adjusting learning rate of group 0 to 2.0000e-06.  
Average epoch loss: 0.5764014214154117  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.27it/s, pos cos\_sim: 0.664  
neg cos\_sim: 0.896]

Average cos sim (pos, neg): 0.785200, 0.840454  
Epoch: 11 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 163.86it/s]

Generated data n=7693

train batches: 100%| | 121/121 [01:44<00:00, 1.16it/s, loss: 0.251]

Adjusting learning rate of group 0 to 2.0000e-06.  
Average epoch loss: 0.53346385880689  
Saved model to: /Users/johnsonzhou/git/comp90042-project/result/models/model\_03a\_bert\_base\_cos\_sim\_2023\_05\_03\_20\_43.pth

dev batches: 100%| | 32/32 [00:07<00:00, 4.32it/s, pos cos\_sim: 0.635  
neg cos\_sim: 0.887]

Average cos sim (pos, neg): 0.762964, 0.835725  
Epoch: 12 of 100

Generate siamese dataset

claims: 100%| | 1228/1228 [00:07<00:00, 174.40it/s]

Generated data n=7693

train batches: 14%| | 17/121 [00:15<01:26, 1.20it/s, loss: 0.528]