

model_06_bert_cross_encoder_classification_grid_search_2

May 8, 2023

1 Model 06 Bert Cross Entropy Classification for Label Prediction

Prediction of claim labels based on the matched evidence.

1.1 Setup

1.1.1 Working Directory

```
[ ]: # Change the working directory to project root
from pathlib import Path
import os
ROOT_DIR = Path.cwd()
while not ROOT_DIR.joinpath("src").exists():
    ROOT_DIR = ROOT_DIR.parent
os.chdir(ROOT_DIR)
```

1.1.2 Dependencies

```
[ ]: # Imports and dependencies
import torch
from torch.utils.data import DataLoader
from torch.nn import CrossEntropyLoss
from torch.optim import AdamW
from torch.optim.lr_scheduler import LinearLR
from torcheval.metrics import MulticlassAccuracy, MulticlassF1Score

from src.logger import SimpleLogger
from src.model_05 import BertCrossEncoderClassifier
from src.data import LabelClassificationDataset
from src.torch_utils import get_torch_device
import json
from dataclasses import dataclass
from typing import List, Union, Tuple
from tqdm import tqdm
import random
import numpy as np
from datetime import datetime
from math import exp
```

```
from sklearn.model_selection import ParameterGrid

TORCH_DEVICE = get_torch_device()
```

/opt/homebrew/Caskroom/miniconda/base/envs/comp90042_project/lib/python3.8/site-packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html

```
from .autonotebook import tqdm as notebook_tqdm
```

Torch device is 'mps'

1.1.3 File paths

```
[ ]: MODEL_PATH = ROOT_DIR.joinpath("./result/models/*")
DATA_PATH = ROOT_DIR.joinpath("./data/*")
LOG_PATH = ROOT_DIR.joinpath("./result/logs/*")
SHORTLIST_PATH = ROOT_DIR.joinpath("./result/pipeline/shortlisting_v2/*")

run_time = datetime.now().strftime('%Y_%m_%d_%H_%M')
```

1.2 Training Loop

```
[ ]: def training_loop(
    model,
    claims_paths:List[Path],
    save_path:Path=None,
    warmup:float=0.1,
    lr:float=0.00005, # 5e-5
    weight_decay:float=0.01,
    normalize_text:bool=True,
    max_length:int=128,
    dropout:float=None,
    n_epochs:int=5,
    batch_size:int=64,
):
    # Generate training dataset
    train_data = LabelClassificationDataset(
        claims_paths=claims_paths,
        training=True,
    )
    train_dataloader = DataLoader(
        dataset=train_data,
        shuffle=True,
        batch_size=batch_size
    )
```

```

# Generate evaluation dataset
dev_data = LabelClassificationDataset(
    claims_paths=[Path("./data/dev-claims.json")],
    training=True,
)
dev_dataloader = DataLoader(
    dataset=dev_data,
    shuffle=False,
    batch_size=batch_size
)

# Loss function
loss_fn = CrossEntropyLoss()

# Optimizer
optimizer = AdamW(
    params=model.parameters(),
    lr=lr,
    weight_decay=weight_decay
)

# Scheduler
scheduler = LinearLR(
    optimizer=optimizer,
    total_iters=warmup * len(train_dataloader),
    verbose=False
)

# Metrics
accuracy_fn = MulticlassAccuracy()
f1_fn = MulticlassF1Score()

# Training epochs -----

best_epoch_loss = 999
best_epoch_f1 = -1
best_epoch_acc = -1
best_epoch = 0
for epoch in range(n_epochs):

    print(f"Epoch: {epoch + 1} of {n_epochs}\n")

    # Run training -----
    model.train()

    train_batches = tqdm(train_dataloader, desc="train batches")
    running_losses = []

```

```

for batch in train_batches:
    claim_texts, evidence_texts, labels, claim_ids, evidence_ids = batch
    texts = list(zip(claim_texts, evidence_texts))

    # Reset optimizer
    optimizer.zero_grad()

    # Forward + loss
    output, logits, seq = model(
        texts=texts,
        normalize_text=normalize_text,
        max_length=max_length,
        dropout=dropout
    )
    loss = loss_fn(logits, labels)

    # Backward + optimizer
    loss.backward()
    optimizer.step()

    # Update running loss
    batch_loss = loss.item() * len(batch)
    running_losses.append(batch_loss)

    train_batches.postfix = f"loss: {batch_loss:.3f}"

    # Update scheduler
    scheduler.step()

    continue

# Epoch loss
epoch_loss = np.average(running_losses)
print(f"Average epoch loss: {epoch_loss:.3f}")

# Run evaluation -----
model.eval()

dev_batches = tqdm(dev_dataloader, desc="dev batches")
dev_acc = []
dev_f1 = []
for batch in dev_batches:
    claim_texts, evidence_texts, labels, claim_ids, evidence_ids = batch
    texts = list(zip(claim_texts, evidence_texts))

    # Forward
    output, logits, seq = model(

```

```

        texts=texts,
        normalize_text=normalize_text,
        max_length=max_length,
        dropout=dropout
    )

    # Prediction
    predicted = torch.argmax(output, dim=1)

    # Metrics
    accuracy_fn.update(predicted.cpu(), labels.cpu())
    f1_fn.update(predicted.cpu(), labels.cpu())

    acc = accuracy_fn.compute()
    f1 = f1_fn.compute()

    dev_acc.append(acc)
    dev_f1.append(f1)

    dev_batches.postfix = f" acc: {acc:.3f}, f1: {f1:.3f}"

    continue

# Consider metrics
    epoch_acc = np.average(dev_acc)
    print(f"Average epoch accuracy: {epoch_acc:.3f}")

    epoch_f1 = np.average(dev_f1)
    print(f"Average epoch f1: {epoch_f1:.3f}")

    if epoch_acc > best_epoch_acc:
        best_epoch_acc = epoch_acc

    if epoch_f1 > best_epoch_f1:
        best_epoch_f1 = epoch_f1
        best_epoch = epoch + 1

# Save model -----

# Save the model with the best f1 score
    if save_path and epoch_f1 >= best_epoch_f1:
        torch.save(model, save_path)
        print(f"Saved model to: {save_path}")

print("Done!")
return best_epoch_acc, best_epoch_f1, best_epoch

```

1.3 Load model

Use a blank pre-trained

```
[ ]: model = BertCrossEncoderClassifier(
    pretrained_name="bert-base-uncased",
    n_classes=3,
    device=TORCH_DEVICE
)
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertModel: ['cls.seq_relationship.bias',

'cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.weight',
'cls.predictions.decoder.weight', 'cls.predictions.transform.dense.weight',
'cls.predictions.transform.LayerNorm.weight',
'cls.predictions.transform.dense.bias', 'cls.predictions.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Or load one previously trained

```
[ ]: # MODEL_SAVE_PATH = MODEL_PATH.with_name("")
# with open(MODEL_PATH.with_name(MODEL_SAVE_PATH), mode="rb") as f:
#     model = torch.load(f, map_location=TORCH_DEVICE)
```

1.4 Training and evaluation loop

```
[ ]: # training_loop(
#     model=model,
#     claims_paths=[
#         DATA_PATH.with_name("train-claims.json")
#     ],
#     save_path=MODEL_PATH.
#         with_name(f"model_06_bert_cross_encoder_label_{run_time}.pth"),
#     warmup=0.1,
#     lr=0.00005, # 5e-5
#     weight_decay=0.01,
#     normalize_text=True,
#     max_length=512,
#     dropout=None,
#     n_epochs=1,
#     batch_size=24,
# )
```

1.5 Tune hyperparameters

```
[ ]: hyperparams = ParameterGrid(param_grid={
    "claims_paths": [
        DATA_PATH.with_name("train-claims.json")
    ],
    "warmup": [0.1],
    "lr": [0.000005],
    "weight_decay": [0.02],
    "normalize_text": [True],
    "max_length": [512],
    "dropout": [0.1, 0.2],
    "n_epochs": [3],
    "batch_size": [8, 32],
    "freeze_bert": [False]
})

[ ]: import warnings
warnings.filterwarnings('ignore')

[ ]: with SimpleLogger("model_06_cross_encoder_retrieval") as logger:
    logger.set_stream_handler()
    logger.set_file_handler(
        log_path=LOG_PATH,
        filename="model_06_hyperparam_tuning.txt"
    )
    best_f1 = -1
    best_params = {}
    for hyperparam in hyperparams:
        model = BertCrossEncoderClassifier(
            pretrained_name="bert-base-uncased",
            n_classes=3,
            device=TORCH_DEVICE
        )

        model_param = hyperparam.copy()

        # Freeze bert parameters if desired
        if "freeze_bert" in model_param.keys():
            if hyperparam["freeze_bert"] is True:
                for param in model.bert.parameters():
                    param.requires_grad = False
                del model_param["freeze_bert"]

        logger.info("\n== RUN")
        logger.info(hyperparam)
```

```

        accuracy, f1, epoch = training_loop(model=model, **model_param)

        logger.info(f"run_best_epoch: {epoch}, run_best_acc: {accuracy},
↪run_best_f1: {f1}")

        if f1 > best_f1:
            best_f1 = f1
            best_params = hyperparam

        logger.info(f"\n== CURRENT BEST F1: {best_f1}")
        logger.info(best_params)

```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertModel: ['cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.weight', 'cls.predictions.decoder.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

2023-05-08 13:47:33 model_06_cross_encoder_retrieval:INFO

== RUN

2023-05-08 13:47:33 model_06_cross_encoder_retrieval:INFO

```

{'batch_size': 8, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.1, 'freeze_bert': False, 'lr': 5e-06, 'max_length': 512,
'n_epochs': 3, 'normalize_text': True, 'warmup': 0.1, 'weight_decay': 0.02}

```

Torch device is 'mps'

claims: 100%| | 1228/1228 [00:00<00:00, 475411.23it/s]

generated dataset n=3730

Torch device is 'mps'

claims: 100%| | 154/154 [00:00<00:00, 544622.95it/s]

generated dataset n=433

Epoch: 1 of 3

train batches: 100%| | 467/467 [03:47<00:00, 2.05it/s, loss: 1.060]

Average epoch loss: 4.479

dev batches: 100%| | 55/55 [00:07<00:00, 7.35it/s, acc: 0.564, f1: 0.564]

Average epoch accuracy: 0.601

Average epoch f1: 0.601

Epoch: 2 of 3

train batches: 100%| | 467/467 [03:47<00:00, 2.06it/s, loss: 1.061]

Average epoch loss: 3.305

dev batches: 100%| | 55/55 [00:07<00:00, 7.50it/s, acc: 0.558, f1: 0.558]

Average epoch accuracy: 0.569

Average epoch f1: 0.569

Epoch: 3 of 3

train batches: 100%| | 467/467 [03:46<00:00, 2.06it/s, loss: 0.201]

Average epoch loss: 2.120

dev batches: 100%| | 55/55 [00:07<00:00, 7.57it/s, acc: 0.570, f1: 0.570]

Average epoch accuracy: 0.570

Average epoch f1: 0.570

Done!

2023-05-08 13:59:18 model_06_cross_encoder_retrieval:INFO

run_best_epoch: 1, run_best_acc: 0.601498007774353, run_best_f1: 0.601498007774353

2023-05-08 13:59:18 model_06_cross_encoder_retrieval:INFO

== CURRENT BEST F1: 0.601498007774353

2023-05-08 13:59:18 model_06_cross_encoder_retrieval:INFO

```
{'batch_size': 8, 'claims_paths':  
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],  
'dropout': 0.1, 'freeze_bert': False, 'lr': 5e-06, 'max_length': 512,  
'n_epochs': 3, 'normalize_text': True, 'warmup': 0.1, 'weight_decay': 0.02}
```

Some weights of the model checkpoint at bert-base-uncased were not used when

initializing BertModel: ['cls.seq_relationship.bias',
'cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.weight',
'cls.predictions.decoder.weight', 'cls.predictions.transform.dense.weight',
'cls.predictions.transform.LayerNorm.weight',
'cls.predictions.transform.dense.bias', 'cls.predictions.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a

BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

2023-05-08 13:59:20 model_06_cross_encoder_retrieval:INFO

== RUN

2023-05-08 13:59:20 model_06_cross_encoder_retrieval:INFO

```
{'batch_size': 8, 'claims_paths':  
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],  
'dropout': 0.2, 'freeze_bert': False, 'lr': 5e-06, 'max_length': 512,  
'n_epochs': 3, 'normalize_text': True, 'warmup': 0.1, 'weight_decay': 0.02}
```

Torch device is 'mps'

claims: 100%| | 1228/1228 [00:00<00:00, 441392.18it/s]

generated dataset n=3730

Torch device is 'mps'

claims: 100%| | 154/154 [00:00<00:00, 618700.02it/s]

generated dataset n=433

Epoch: 1 of 3

train batches: 100%| | 467/467 [03:47<00:00, 2.05it/s, loss: 4.158]

Average epoch loss: 4.537

dev batches: 100%| | 55/55 [00:07<00:00, 7.55it/s, acc: 0.561, f1:
0.561]

Average epoch accuracy: 0.616

Average epoch f1: 0.616

Epoch: 2 of 3

train batches: 100%| | 467/467 [03:48<00:00, 2.05it/s, loss: 3.084]

Average epoch loss: 3.305

dev batches: 100%| | 55/55 [00:07<00:00, 7.56it/s, acc: 0.581, f1:
0.581]

Average epoch accuracy: 0.585

Average epoch f1: 0.585

Epoch: 3 of 3

train batches: 100%| | 467/467 [03:48<00:00, 2.05it/s, loss: 4.307]

Average epoch loss: 2.327

dev batches: 100%| | 55/55 [00:07<00:00, 7.54it/s, acc: 0.587, f1: 0.587]

Average epoch accuracy: 0.591

Average epoch f1: 0.591

Done!

2023-05-08 14:11:07 model_06_cross_encoder_retrieval:INFO

run_best_epoch: 1, run_best_acc: 0.6160757541656494, run_best_f1: 0.6160757541656494

2023-05-08 14:11:07 model_06_cross_encoder_retrieval:INFO

== CURRENT BEST F1: 0.6160757541656494

2023-05-08 14:11:07 model_06_cross_encoder_retrieval:INFO

{'batch_size': 8, 'claims_paths': [PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')], 'dropout': 0.2, 'freeze_bert': False, 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text': True, 'warmup': 0.1, 'weight_decay': 0.02}

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertModel: ['cls.seq_relationship.bias',

'cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.weight', 'cls.predictions.decoder.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

2023-05-08 14:11:08 model_06_cross_encoder_retrieval:INFO

== RUN

2023-05-08 14:11:08 model_06_cross_encoder_retrieval:INFO

{'batch_size': 32, 'claims_paths': [PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')], 'dropout': 0.1, 'freeze_bert': False, 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text': True, 'warmup': 0.1, 'weight_decay': 0.02}

Torch device is 'mps'

claims: 100%| | 1228/1228 [00:00<00:00, 631046.96it/s]

generated dataset n=3730

Torch device is 'mps'

claims: 100%| | 154/154 [00:00<00:00, 621079.63it/s]

generated dataset n=433

Epoch: 1 of 3

train batches: 100%| | 117/117 [03:22<00:00, 1.73s/it, loss: 4.321]

Average epoch loss: 4.887

dev batches: 100%| | 14/14 [00:07<00:00, 1.97it/s, acc: 0.527, f1: 0.527]

Average epoch accuracy: 0.580

Average epoch f1: 0.580

Epoch: 2 of 3

train batches: 100%| | 117/117 [03:22<00:00, 1.73s/it, loss: 2.214]

Average epoch loss: 4.020

dev batches: 100%| | 14/14 [00:06<00:00, 2.00it/s, acc: 0.527, f1: 0.527]

Average epoch accuracy: 0.546

Average epoch f1: 0.546

Epoch: 3 of 3

train batches: 100%| | 117/117 [03:21<00:00, 1.72s/it, loss: 3.609]

Average epoch loss: 3.360

dev batches: 100%| | 14/14 [00:06<00:00, 2.00it/s, acc: 0.530, f1: 0.530]

Average epoch accuracy: 0.539

Average epoch f1: 0.539

Done!

2023-05-08 14:21:37 model_06_cross_encoder_retrieval:INFO

run_best_epoch: 1, run_best_acc: 0.5803744196891785, run_best_f1: 0.5803744196891785

2023-05-08 14:21:37 model_06_cross_encoder_retrieval:INFO

== CURRENT BEST F1: 0.6160757541656494

2023-05-08 14:21:37 model_06_cross_encoder_retrieval:INFO

{'batch_size': 8, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.2, 'freeze_bert': False, 'lr': 5e-06, 'max_length': 512,
'n_epochs': 3, 'normalize_text': True, 'warmup': 0.1, 'weight_decay': 0.02}

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertModel: ['cls.seq_relationship.bias', 'cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.weight', 'cls.predictions.decoder.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

2023-05-08 14:21:39 model_06_cross_encoder_retrieval:INFO

== RUN

2023-05-08 14:21:39 model_06_cross_encoder_retrieval:INFO
{'batch_size': 32, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.2, 'freeze_bert': False, 'lr': 5e-06, 'max_length': 512,
'n_epochs': 3, 'normalize_text': True, 'warmup': 0.1, 'weight_decay': 0.02}

Torch device is 'mps'

claims: 100%| | 1228/1228 [00:00<00:00, 601635.94it/s]

generated dataset n=3730

Torch device is 'mps'

claims: 100%| | 154/154 [00:00<00:00, 628941.40it/s]

generated dataset n=433

Epoch: 1 of 3

train batches: 100%| | 117/117 [03:23<00:00, 1.74s/it, loss: 4.077]

Average epoch loss: 4.923

dev batches: 100%| | 14/14 [00:06<00:00, 2.00it/s, acc: 0.439, f1: 0.439]

Average epoch accuracy: 0.450

Average epoch f1: 0.450

Epoch: 2 of 3

train batches: 100%| | 117/117 [03:21<00:00, 1.72s/it, loss: 2.892]

Average epoch loss: 4.342

dev batches: 100%| | 14/14 [00:06<00:00, 2.01it/s, acc: 0.498, f1: 0.498]

Average epoch accuracy: 0.491

Average epoch f1: 0.491

Epoch: 3 of 3

train batches: 100%| | 117/117 [03:21<00:00, 1.72s/it, loss: 2.412]

Average epoch loss: 3.807

dev batches: 100%| | 14/14 [00:06<00:00, 2.00it/s, acc: 0.517, f1: 0.517]

Average epoch accuracy: 0.518

Average epoch f1: 0.518

Done!

2023-05-08 14:32:07 model_06_cross_encoder_retrieval:INFO

run_best_epoch: 3, run_best_acc: 0.5175042748451233, run_best_f1: 0.5175042748451233

2023-05-08 14:32:07 model_06_cross_encoder_retrieval:INFO

== CURRENT BEST F1: 0.6160757541656494

2023-05-08 14:32:07 model_06_cross_encoder_retrieval:INFO

{'batch_size': 8, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.2, 'freeze_bert': False, 'lr': 5e-06, 'max_length': 512,
'n_epochs': 3, 'normalize_text': True, 'warmup': 0.1, 'weight_decay': 0.02}