

model_06_bert_cross_encoder_classification_weights_smoothing

May 9, 2023

1 Model 06 Bert Cross Entropy Classification for Label Prediction

Prediction of claim labels based on the matched evidence.

1.1 Setup

1.1.1 Working Directory

```
[ ]: # Change the working directory to project root
from pathlib import Path
import os
ROOT_DIR = Path.cwd()
while not ROOT_DIR.joinpath("src").exists():
    ROOT_DIR = ROOT_DIR.parent
os.chdir(ROOT_DIR)
```

1.1.2 Dependencies

```
[ ]: # Imports and dependencies
import torch
from torch.utils.data import DataLoader
from torch.nn import CrossEntropyLoss
from torch.optim import AdamW
from torch.optim.lr_scheduler import LinearLR
from torcheval.metrics import MulticlassAccuracy, MulticlassF1Score

from src.logger import SimpleLogger
from src.model_05 import BertCrossEncoderClassifier
from src.data import LabelClassificationDataset
from src.torch_utils import get_torch_device
import json
from dataclasses import dataclass
from typing import List, Union, Tuple
from tqdm import tqdm
import random
import numpy as np
from datetime import datetime
from math import exp
```

```
from sklearn.model_selection import ParameterGrid

TORCH_DEVICE = get_torch_device()
```

Torch device is 'mps'

1.1.3 File paths

```
[ ]: MODEL_PATH = ROOT_DIR.joinpath("./result/models/*")
DATA_PATH = ROOT_DIR.joinpath("./data/*")
LOG_PATH = ROOT_DIR.joinpath("./result/logs/*")
SHORTLIST_PATH = ROOT_DIR.joinpath("./result/pipeline/shortlisting_v2/*")

run_time = datetime.now().strftime('%Y_%m_%d_%H_%M')
```

1.2 Training Loop

```
[ ]: def training_loop(
    model,
    claims_paths:List[Path],
    save_path:Path=None,
    label_weight:list=None,
    label_smoothing:float=None,
    warmup:float=0.1,
    lr:float=0.00005, # 5e-5
    weight_decay:float=0.01,
    normalize_text:bool=True,
    max_length:int=128,
    dropout:float=None,
    n_epochs:int=5,
    batch_size:int=64,
):
    # Generate training dataset
    train_data = LabelClassificationDataset(
        claims_paths=claims_paths,
        training=True,
    )
    train_dataloader = DataLoader(
        dataset=train_data,
        shuffle=True,
        batch_size=batch_size
    )

    # Generate evaluation dataset
    dev_data = LabelClassificationDataset(
        claims_paths=[Path("./data/dev-claims.json")],
        training=True,
```

```

)
dev_dataloader = DataLoader(
    dataset=dev_data,
    shuffle=False,
    batch_size=batch_size
)

# Loss function
loss_fn = CrossEntropyLoss(
    weight=torch.tensor(label_weight, device=TORCH_DEVICE),
    label_smoothing=label_smoothing
)

# Optimizer
optimizer = AdamW(
    params=model.parameters(),
    lr=lr,
    weight_decay=weight_decay
)

# Scheduler
scheduler = LinearLR(
    optimizer=optimizer,
    total_iters=warmup * len(train_dataloader),
    verbose=False
)

# Metrics
accuracy_fn = MulticlassAccuracy()
f1_fn = MulticlassF1Score()

# Training epochs -----

best_epoch_loss = 999
best_epoch_f1 = -1
best_epoch_acc = -1
best_epoch = 0
for epoch in range(n_epochs):

    print(f"\nEpoch: {epoch + 1} of {n_epochs}\n")

    # Run training -----
    model.train()

    train_batches = tqdm(train_dataloader, desc="train batches")
    running_losses = []
    for batch in train_batches:

```

```

claim_texts, evidence_texts, labels, claim_ids, evidence_ids = batch
texts = list(zip(claim_texts, evidence_texts))

# Reset optimizer
optimizer.zero_grad()

# Forward + loss
output, logits, seq = model(
    texts=texts,
    normalize_text=normalize_text,
    max_length=max_length,
    dropout=dropout
)
loss = loss_fn(logits, labels)

# Backward + optimizer
loss.backward()
optimizer.step()

# Update running loss
batch_loss = loss.item() * len(batch)
running_losses.append(batch_loss)

train_batches.postfix = f"loss: {batch_loss:.3f}"

# Update scheduler
scheduler.step()

continue

# Epoch loss
epoch_loss = np.average(running_losses)
print(f"Average epoch loss: {epoch_loss:.3f}")

# Run evaluation -----
model.eval()

dev_batches = tqdm(dev_dataloader, desc="dev batches")
dev_acc = []
dev_f1 = []
for batch in dev_batches:
    claim_texts, evidence_texts, labels, claim_ids, evidence_ids = batch
    texts = list(zip(claim_texts, evidence_texts))

    # Forward
    output, logits, seq = model(
        texts=texts,

```

```

        normalize_text=normalize_text,
        max_length=max_length,
        dropout=dropout
    )

    # Prediction
    predicted = torch.argmax(output, dim=1)

    # Metrics
    accuracy_fn.update(predicted.cpu(), labels.cpu())
    f1_fn.update(predicted.cpu(), labels.cpu())

    acc = accuracy_fn.compute()
    f1 = f1_fn.compute()

    dev_acc.append(acc)
    dev_f1.append(f1)

    dev_batches.postfix = f" acc: {acc:.3f}, f1: {f1:.3f}"

    continue

# Consider metrics
    epoch_acc = np.average(dev_acc)
    print(f"Average epoch accuracy: {epoch_acc:.3f}")

    epoch_f1 = np.average(dev_f1)
    print(f"Average epoch f1: {epoch_f1:.3f}")

    if epoch_acc > best_epoch_acc:
        best_epoch_acc = epoch_acc

    if epoch_f1 > best_epoch_f1:
        best_epoch_f1 = epoch_f1
        best_epoch = epoch + 1

# Save model -----

# Save the model with the best f1 score
    if save_path and epoch_f1 >= best_epoch_f1:
        torch.save(model, save_path)
        print(f"Saved model to: {save_path}")

print("Done!")
return best_epoch_acc, best_epoch_f1, best_epoch

```

1.3 Load model

Use a blank pre-trained

```
[ ]: # model = BertCrossEncoderClassifier(  
#     pretrained_name="bert-base-uncased",  
#     n_classes=3,  
#     device=TORCH_DEVICE  
# )
```

Or load one previously trained

```
[ ]: # MODEL_SAVE_PATH = MODEL_PATH.with_name("")  
# with open(MODEL_PATH.with_name(MODEL_SAVE_PATH), mode="rb") as f:  
#     model = torch.load(f, map_location=TORCH_DEVICE)
```

1.4 Training and evaluation loop

```
[ ]: # training_loop(  
#     model=model,  
#     claims_paths=[  
#         DATA_PATH.with_name("train-claims.json")  
#     ],  
#     save_path=MODEL_PATH.  
#         ↪with_name(f"model_06_bert_base_uncased_cross_encoder_label_{run_time}.pth"),  
#     warmup=0.1,  
#     lr=0.000005, # 5e-6  
#     weight_decay=0.02,  
#     normalize_text=True,  
#     max_length=512,  
#     dropout=0.1,  
#     n_epochs=1,  
#     batch_size=24,  
# )
```

1.5 Tune hyperparameters

```
[ ]: hyperparams = ParameterGrid(param_grid={  
    "claims_paths": [[  
        DATA_PATH.with_name("train-claims.json")  
    ]],  
    "warmup": [0.1],  
    "lr": [0.000005],  
    "weight_decay": [0.02],  
    "normalize_text": [True],  
    "max_length": [512],  
    "dropout": [0.1],
```

```

    "n_epochs": [3],
    "batch_size": [24],
    "freeze_bert": [False],
    "label_weight": [
        [2, 1.2, 1],
        [1, 0.6, 0.4],
    ],
    "label_smoothing": [0.0, 0.2]
})

```

```

[ ]: import warnings
warnings.filterwarnings('ignore')

```

```

[ ]: with SimpleLogger("model_06_cross_encoder_retrieval") as logger:
    logger.set_stream_handler()
    logger.set_file_handler(
        log_path=LOG_PATH,
        filename="model_06_hyperparam_tuning.txt"
    )
    best_f1 = -1
    best_params = {}
    for hyperparam in hyperparams:
        model = BertCrossEncoderClassifier(
            pretrained_name="bert-base-uncased",
            n_classes=3,
            device=TORCH_DEVICE
        )

        model_param = hyperparam.copy()

        # Freeze bert parameters if desired
        if "freeze_bert" in model_param.keys():
            if hyperparam["freeze_bert"] is True:
                for param in model.bert.parameters():
                    param.requires_grad = False
                del model_param["freeze_bert"]

        logger.info("\n== RUN")
        logger.info(hyperparam)

        accuracy, f1, epoch = training_loop(model=model, **model_param)

        logger.info(f"run_best_epoch: {epoch}, run_best_acc: {accuracy},
↪run_best_f1: {f1}")

        if f1 > best_f1:
            best_f1 = f1

```

```

best_params = hyperparam

logger.info(f"\n== CURRENT BEST F1: {best_f1}")
logger.info(best_params)

```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.bias', 'cls.seq_relationship.bias', 'cls.predictions.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight', 'cls.seq_relationship.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

2023-05-09 09:42:38 model_06_cross_encoder_retrieval:INFO

== RUN

2023-05-09 09:42:38 model_06_cross_encoder_retrieval:INFO

```

{'batch_size': 24, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.1, 'freeze_bert': False, 'label_smoothing': 0.0, 'label_weight':
[2, 1.2, 1], 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text':
True, 'warmup': 0.1, 'weight_decay': 0.02}

```

Torch device is 'mps'

claims: 100%| | 1228/1228 [00:00<00:00, 821861.39it/s]

generated dataset n=3730

Torch device is 'mps'

claims: 100%| | 154/154 [00:00<00:00, 367628.24it/s]

generated dataset n=433

Epoch: 1 of 3

train batches: 100%| | 156/156 [03:24<00:00, 1.31s/it, loss: 4.121]

Average epoch loss: 5.123

dev batches: 100%| | 19/19 [00:07<00:00, 2.66it/s, acc: 0.566, f1: 0.566]

Average epoch accuracy: 0.603
Average epoch f1: 0.603

Epoch: 2 of 3

train batches: 100%| | 156/156 [03:23<00:00, 1.30s/it, loss: 2.406]

Average epoch loss: 4.072

dev batches: 100%| | 19/19 [00:07<00:00, 2.63it/s, acc: 0.569, f1:
0.569]

Average epoch accuracy: 0.584
Average epoch f1: 0.584

Epoch: 3 of 3

train batches: 100%| | 156/156 [03:23<00:00, 1.30s/it, loss: 3.258]

Average epoch loss: 3.357

dev batches: 100%| | 19/19 [00:07<00:00, 2.61it/s, acc: 0.583, f1:
0.583]

Average epoch accuracy: 0.583
Average epoch f1: 0.583

Done!

2023-05-09 09:53:11 model_06_cross_encoder_retrieval:INFO
run_best_epoch: 1, run_best_acc: 0.6030740737915039, run_best_f1:
0.6030740737915039

2023-05-09 09:53:11 model_06_cross_encoder_retrieval:INFO

== CURRENT BEST F1: 0.6030740737915039

2023-05-09 09:53:11 model_06_cross_encoder_retrieval:INFO
{'batch_size': 24, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.1, 'freeze_bert': False, 'label_smoothing': 0.0, 'label_weight':
[2, 1.2, 1], 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text':
True, 'warmup': 0.1, 'weight_decay': 0.02}

Some weights of the model checkpoint at bert-base-uncased were not used when
initializing BertModel: ['cls.predictions.transform.LayerNorm.bias',
'cls.seq_relationship.bias', 'cls.predictions.bias',
'cls.predictions.transform.LayerNorm.weight',
'cls.predictions.transform.dense.bias',
'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight',

```
'cls.seq_relationship.weight']
- This IS expected if you are initializing BertModel from the checkpoint of a
model trained on another task or with another architecture (e.g. initializing a
BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of
a model that you expect to be exactly identical (initializing a
BertForSequenceClassification model from a BertForSequenceClassification model).
2023-05-09 09:53:13 model_06_cross_encoder_retrieval:INFO
```

== RUN

```
2023-05-09 09:53:13 model_06_cross_encoder_retrieval:INFO
{'batch_size': 24, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.1, 'freeze_bert': False, 'label_smoothing': 0.0, 'label_weight':
[1, 0.6, 0.4], 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text':
True, 'warmup': 0.1, 'weight_decay': 0.02}
```

Torch device is 'mps'

```
claims: 100%|          | 1228/1228 [00:00<00:00, 576775.51it/s]
```

generated dataset n=3730

Torch device is 'mps'

```
claims: 100%|          | 154/154 [00:00<00:00, 536480.74it/s]
```

generated dataset n=433

Epoch: 1 of 3

```
train batches: 100%|          | 156/156 [03:25<00:00, 1.32s/it, loss: 4.359]
```

Average epoch loss: 5.414

```
dev batches: 100%|          | 19/19 [00:07<00:00, 2.65it/s, acc: 0.448, f1:
0.448]
```

Average epoch accuracy: 0.445

Average epoch f1: 0.445

Epoch: 2 of 3

```
train batches: 100%|          | 156/156 [03:23<00:00, 1.31s/it, loss: 4.034]
```

Average epoch loss: 4.369

```
dev batches: 100%|          | 19/19 [00:07<00:00, 2.63it/s, acc: 0.482, f1:
0.482]
```

Average epoch accuracy: 0.472
Average epoch f1: 0.472

Epoch: 3 of 3

train batches: 100%| | 156/156 [03:24<00:00, 1.31s/it, loss: 5.414]

Average epoch loss: 3.451

dev batches: 100%| | 19/19 [00:07<00:00, 2.61it/s, acc: 0.498, f1:
0.498]

Average epoch accuracy: 0.497

Average epoch f1: 0.497

Done!

2023-05-09 10:03:49 model_06_cross_encoder_retrieval:INFO

run_best_epoch: 3, run_best_acc: 0.496741384267807, run_best_f1:
0.496741384267807

2023-05-09 10:03:49 model_06_cross_encoder_retrieval:INFO

== CURRENT BEST F1: 0.6030740737915039

2023-05-09 10:03:49 model_06_cross_encoder_retrieval:INFO

{'batch_size': 24, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.1, 'freeze_bert': False, 'label_smoothing': 0.0, 'label_weight':
[2, 1.2, 1], 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text':
True, 'warmup': 0.1, 'weight_decay': 0.02}

Some weights of the model checkpoint at bert-base-uncased were not used when
initializing BertModel: ['cls.predictions.transform.LayerNorm.bias',

'cls.seq_relationship.bias', 'cls.predictions.bias',
'cls.predictions.transform.LayerNorm.weight',
'cls.predictions.transform.dense.bias',
'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight',
'cls.seq_relationship.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a
model trained on another task or with another architecture (e.g. initializing a
BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of
a model that you expect to be exactly identical (initializing a
BertForSequenceClassification model from a BertForSequenceClassification model).

2023-05-09 10:03:51 model_06_cross_encoder_retrieval:INFO

== RUN

```
2023-05-09 10:03:51 model_06_cross_encoder_retrieval:INFO
{'batch_size': 24, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.1, 'freeze_bert': False, 'label_smoothing': 0.2, 'label_weight':
[2, 1.2, 1], 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text':
True, 'warmup': 0.1, 'weight_decay': 0.02}
```

Torch device is 'mps'

```
claims: 100%|      | 1228/1228 [00:00<00:00, 623333.57it/s]
```

generated dataset n=3730

Torch device is 'mps'

```
claims: 100%|      | 154/154 [00:00<00:00, 657093.40it/s]
```

generated dataset n=433

Epoch: 1 of 3

```
train batches: 100%|      | 156/156 [03:25<00:00, 1.31s/it, loss: 4.584]
```

Average epoch loss: 5.382

```
dev batches: 100%|      | 19/19 [00:07<00:00, 2.64it/s, acc: 0.533, f1:
0.533]
```

Average epoch accuracy: 0.573

Average epoch f1: 0.573

Epoch: 2 of 3

```
train batches: 100%|      | 156/156 [03:24<00:00, 1.31s/it, loss: 4.013]
```

Average epoch loss: 4.851

```
dev batches: 100%|      | 19/19 [00:07<00:00, 2.65it/s, acc: 0.548, f1:
0.548]
```

Average epoch accuracy: 0.558

Average epoch f1: 0.558

Epoch: 3 of 3

```
train batches: 100%|      | 156/156 [03:23<00:00, 1.31s/it, loss: 3.575]
```

Average epoch loss: 4.418

```
dev batches: 100%|      | 19/19 [00:07<00:00, 2.65it/s, acc: 0.550, f1:
0.550]
```

Average epoch accuracy: 0.557

Average epoch f1: 0.557

Done!

```
2023-05-09 10:14:27 model_06_cross_encoder_retrieval:INFO
run_best_epoch: 1, run_best_acc: 0.5727986693382263, run_best_f1:
0.5727986693382263
```

```
2023-05-09 10:14:27 model_06_cross_encoder_retrieval:INFO
```

== CURRENT BEST F1: 0.6030740737915039

```
2023-05-09 10:14:27 model_06_cross_encoder_retrieval:INFO
{'batch_size': 24, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.1, 'freeze_bert': False, 'label_smoothing': 0.0, 'label_weight':
[2, 1.2, 1], 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text':
True, 'warmup': 0.1, 'weight_decay': 0.02}
```

Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.bias',

'cls.seq_relationship.bias', 'cls.predictions.bias',
'cls.predictions.transform.LayerNorm.weight',
'cls.predictions.transform.dense.bias',
'cls.predictions.transform.dense.weight', 'cls.predictions.decoder.weight',
'cls.seq_relationship.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```
2023-05-09 10:14:34 model_06_cross_encoder_retrieval:INFO
```

== RUN

```
2023-05-09 10:14:34 model_06_cross_encoder_retrieval:INFO
{'batch_size': 24, 'claims_paths':
[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],
'dropout': 0.1, 'freeze_bert': False, 'label_smoothing': 0.2, 'label_weight':
[1, 0.6, 0.4], 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text':
True, 'warmup': 0.1, 'weight_decay': 0.02}
```

Torch device is 'mps'

claims: 100%| | 1228/1228 [00:00<00:00, 691078.13it/s]

generated dataset n=3730

Torch device is 'mps'

claims: 100%| | 154/154 [00:00<00:00, 292802.73it/s]

generated dataset n=433

Epoch: 1 of 3

train batches: 100%| | 156/156 [03:25<00:00, 1.32s/it, loss: 4.562]

Average epoch loss: 5.361

dev batches: 100%| | 19/19 [00:07<00:00, 2.64it/s, acc: 0.524, f1: 0.524]

Average epoch accuracy: 0.571

Average epoch f1: 0.571

Epoch: 2 of 3

train batches: 100%| | 156/156 [03:24<00:00, 1.31s/it, loss: 4.070]

Average epoch loss: 4.747

dev batches: 100%| | 19/19 [00:07<00:00, 2.63it/s, acc: 0.542, f1: 0.542]

Average epoch accuracy: 0.547

Average epoch f1: 0.547

Epoch: 3 of 3

train batches: 100%| | 156/156 [03:26<00:00, 1.32s/it, loss: 3.959]

Average epoch loss: 4.408

dev batches: 100%| | 19/19 [00:07<00:00, 2.57it/s, acc: 0.549, f1: 0.549]

Average epoch accuracy: 0.553

Average epoch f1: 0.553

Done!

2023-05-09 10:25:13 model_06_cross_encoder_retrieval:INFO

run_best_epoch: 1, run_best_acc: 0.5706400871276855, run_best_f1: 0.5706400871276855

2023-05-09 10:25:13 model_06_cross_encoder_retrieval:INFO

== CURRENT BEST F1: 0.6030740737915039

2023-05-09 10:25:13 model_06_cross_encoder_retrieval:INFO

{'batch_size': 24, 'claims_paths':

[PosixPath('/Users/johnsonzhou/git/comp90042-project/data/train-claims.json')],

```
'dropout': 0.1, 'freeze_bert': False, 'label_smoothing': 0.0, 'label_weight':  
[2, 1.2, 1], 'lr': 5e-06, 'max_length': 512, 'n_epochs': 3, 'normalize_text':  
True, 'warmup': 0.1, 'weight_decay': 0.02}
```